

## 6) Customer Affinities (Brand/Category) — Decayed shares (Dirichlet-smoothed)

### Goal & decision

Score each customer's **brand/category affinity** for **targeting** (audiences/eligibility) and **slotting** (what to surface first).

### Feature sources

- **events** (views/ATC/purchases): `customer_id`, `brand_id`, `category_id`, `ts`
- **customers** (store outputs): `users.affinities[]`

### Compact formula

**Time-decayed event counts per brand** (half-life  $h$  days; event age  $a_e$  days):

$$n_{c,b} = \sum_{e \in \mathcal{E}_{c,b}} 2^{-a_e/h}$$

**Dirichlet-smoothed affinity** (typ.  $\alpha = 1$ ):

$$\text{affinity}_{c,b} = \frac{n_{c,b} + \alpha}{\sum_{b \in B} n_{c,b} + \alpha |B|}$$

### Full expansion (worked)

**Customer's events and parameters:**  $h = 30$  days,  $\alpha = 1$ ,  $B = \{B1, B2, B3\}$ .

B1 ages =  $\{0, 5, 20\}$ ; B2 age = 10; B3 age = 40.

$$n_{c,B1} = 2^{-0/30} + 2^{-5/30} + 2^{-20/30} = 1 + 2^{-1/6} + 2^{-2/3} \approx 2.521$$

$$n_{c,B2} = 2^{-10/30} = 2^{-1/3} \approx 0.794$$

$$n_{c,B3} = 2^{-40/30} = 2^{-4/3} \approx 0.397$$

$$\text{denom} = \sum_{b \in B} n_{c,b} + \alpha |B| = 2.521 + 0.794 + 0.397 + 3 = \mathbf{6.712}$$

$$\text{affinity}_{c,B1} = \frac{2.521 + 1}{6.712} = \mathbf{0.525}$$

$$\text{affinity}_{c,B2} = \frac{0.794 + 1}{6.712} = \mathbf{0.267}$$

$$\text{affinity}_{c,B3} = \frac{0.397 + 1}{6.712} = \mathbf{0.208}$$

### Result & steering

- Persist **top-K** { `id`, `score`, `kind` } per customer for **audiences** and **slotting**.
- Tune **half-life**  $h$  (30–60 d commonly). Use **brand** and/or **category** tracks.

- If events are **sparse**, fall back to **recent-top** brand/category.

▼ **metrics: affinities\_example (with feature\_sources)** — click to expand

```
metrics:
  affinities_example:
    half_life_days: 30
    alpha_smoothing: 1.0
    brands: [B1, B2, B3]
    customer_events:
      # event ages in days (most recent = 0)
      B1: [0, 5, 20]
      B2: [10]
      B3: [40]
    derived:
      n_counts:
        B1: 2.521 # 1 + 2^(-5/30) + 2^(-20/30)
        B2: 0.794
        B3: 0.397
      denom: 6.712 # sum(n_counts) + alpha*|B|
      affinities:
        B1: 0.525
        B2: 0.267
        B3: 0.208
  feature_sources:
    - mongo.events
    - mongo.customers
```

</details>

## 7) Product Assortment Similarity / Clustering — Co-buy graph + cosine (ANN serving)

### Goal & decision

Provide “**similar products**” for **PDP/PLP** and **clusters** for discovery/bundling, especially for the **long tail**.

### Feature sources

- **orders** → basket co-occurrence pairs
- **product\_graph** (derived) → `co_buy_counts` , `neighbors`
- **products** → `category_id` , `brand_id` , `price_cents` , `attributes` (attribute fallback)
- **Outputs** → `products.cluster_id` , `product_features.similar[]`

### Compact formula

**Co-buy row vector for item  $i$**  (counts to other items):

$$\mathbf{c}_i = (c_{i,1}, \dots, c_{i,n})$$

**Cosine similarity:**

$$\cos(i, j) = \frac{\mathbf{c}_i \cdot \mathbf{c}_j}{\|\mathbf{c}_i\| \|\mathbf{c}_j\|}$$

**(Optional) PMI edge weighting → PPMI**

$$\text{PMI}_{ij} = \log\left(\frac{p(i, j)}{p(i)p(j)}\right), \quad \text{PPMI}_{ij} = \max(0, \text{PMI}_{ij})$$

**Serving:** build **ANN index** over item embeddings or normalized co-buy vectors; return **top-K neighbors**; re-rank with **margin/stock** constraints.

## Full expansion (worked)

**Toy co-buy counts (selected pairs):**

$c_{1,2} = 10, c_{1,3} = 12, c_{1,5} = 3; c_{3,1} = 12, c_{3,2} = 5, c_{3,4} = 6; c_{2,5} = 8$  (others = 0).

**Construct co-buy rows over shared index set  $\{2, 3, 4, 5\}$ :**

$\mathbf{c}_1 = [10, 12, 0, 3], \quad \mathbf{c}_3 = [5, 0, 6, 0]$ .

**Cosine pieces (dot and norms):**

$$\mathbf{c}_1 \cdot \mathbf{c}_3 = 10 \cdot 5 + 12 \cdot 0 + 0 \cdot 6 + 3 \cdot 0 = \mathbf{50}$$

$$\|\mathbf{c}_1\| = \sqrt{10^2 + 12^2 + 0^2 + 3^2} = \sqrt{253} \approx \mathbf{15.90597}$$

$$\|\mathbf{c}_3\| = \sqrt{5^2 + 0^2 + 6^2 + 0^2} = \sqrt{61} \approx \mathbf{7.81025}$$

**Cosine similarity (P1 vs P3):**

$$\cos(1, 3) = \frac{50}{15.90597 \cdot 7.81025} = 0.40248 \approx \mathbf{0.4025}$$

**Neighbor selection & clustering:** repeat vs other items → pick **top-K neighbors**; cluster the neighbor graph via **Louvain/k-means**; write `products.cluster_id`.

## Result & steering

- **PDP neighbors:** return top-K by cosine; re-rank to **penalize low-stock** and **low-margin** items.
- **Assortment:** use `cluster_id` for browse pages, **markdown cross-effects**, and **bundle** candidates.
- **Cold start:** if counts are sparse, fall back to **attribute/semantic** similarity until interactions accrue.

▼ **metrics: product\_similarity\_example (with feature\_sources)** — click to expand

```

metrics:
  product_similarity_example:
    items: [P1, P2, P3, P4, P5]
    co_buy_counts:
      P1: { P2: 10, P3: 12, P4: 0, P5: 3 }
      P3: { P1: 12, P2: 5, P4: 6, P5: 0 }
    vectors:
      P1: [10, 12, 0, 3]      # to [P2,P3,P4,P5]
      P3: [5, 0, 6, 0]
    cosine:
      dot: 50
      norm_P1: 15.90597
      norm_P3: 7.81025
      cos_P1_P3: 0.4025
  feature_sources:
    - mongo.orders
    - mongo.product_graph
    - mongo.products

</details>

```

## 8) PLP/Search Ranking (LTR) — BM25/ANN retrieval → pairwise RankNet

### Goal & decision

Order PLP/Search results to **maximize relevance and conversion (nDCG@K)**, honoring **stock/price/promo freshness**.

### Feature sources

- **search\_logs**: query, impressions, clicks, purchases
- **product\_features**: bm25 , semantic\_ann , margin\_pct , stock\_cover\_days , promo\_flag , price\_norm
- **products** (fresh price/stock)

### Compact formula

**Pairwise RankNet loss** (clicked  $i$ , not-clicked  $j$ ):

$$\mathcal{L}_{ij} = \log(1 + \exp(-(s_i - s_j)))$$

with **score function**:

$$s_k = \mathbf{w}^\top \mathbf{x}_k$$

**Gradients wrt scores**:

$$\frac{\partial \mathcal{L}_{ij}}{\partial s_i} = \sigma(s_j - s_i), \quad \frac{\partial \mathcal{L}_{ij}}{\partial s_j} = -\sigma(s_j - s_i), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

## Full expansion (worked)

One clicked, one not-clicked:  $s_i = 1.20$ ,  $s_j = 0.50 \Rightarrow s_i - s_j = 0.70$ .

Loss:

$$\mathcal{L}_{ij} = \log(1 + e^{-(s_i - s_j)}) = \log(1 + e^{-0.70}) = \log(1 + 0.496585) \approx \mathbf{0.403}$$

Gradient signal:

$$\sigma(s_j - s_i) = \sigma(-0.70) = \frac{1}{1 + e^{0.70}} \approx \mathbf{0.3318}$$

Interpretation: push  $s_i$  **up** by +0.3318 (descent uses  $-\partial\mathcal{L}/\partial s_i$ ); push  $s_j$  **down** by 0.3318.

## Scoring (illustrative)

Let

$$s_k = w_0 + w_1 \text{bm25}_k + w_2 \text{sem}_k + w_3 \text{stockCover}_k + w_4 \text{promo}_k, \quad \mathbf{w} = (0, 0.6, 0.8, 0.05, 0.2).$$

Inputs:

A: bm25 = 1.4, sem = 0.9, stockCover = 9, promo = 1

B: bm25 = 1.2, sem = 0.7, stockCover = 6, promo = 0

C: bm25 = 0.9, sem = 0.5, stockCover = 4, promo = 0

Scores:

$$s_A = 0.6 \cdot 1.4 + 0.8 \cdot 0.9 + 0.05 \cdot 9 + 0.2 \cdot 1 = 0.84 + 0.72 + 0.45 + 0.20 = \mathbf{2.21}$$

$$s_B = 0.6 \cdot 1.2 + 0.8 \cdot 0.7 + 0.05 \cdot 6 + 0.2 \cdot 0 = 0.72 + 0.56 + 0.30 + 0 = \mathbf{1.58}$$

$$s_C = 0.6 \cdot 0.9 + 0.8 \cdot 0.5 + 0.05 \cdot 4 + 0.2 \cdot 0 = 0.54 + 0.40 + 0.20 + 0 = \mathbf{1.14}$$

Ranking:  $A > B > C$ . Apply **OOS/stock/price freshness gates** and **exposure caps after** scoring.

## Result & steering

- Use **BM25** u **ANN** for retrieval; rank with **LTR (RankNet/LambdaMART)**.
- If **price/stock features** are stale (> 5 min), **fall back** to lexical + business rules.

▼ metrics: plp\_ltr\_example (with feature\_sources) — click to expand

```

metrics:
  plp_ltr_example:
    pairwise_loss:
      s_i: 1.20
      s_j: 0.50
      loss_logistic: 0.403
      grad_sigma_sj_minus_si: 0.3318
    scoring_weights:
      w0: 0.0
      w1_bm25: 0.6
      w2_sem: 0.8
      w3_stockCover: 0.05
      w4_promo: 0.2
    items:
      - { id: A, bm25: 1.4, sem: 0.9, stock_cover_days: 9, promo: 1, score: 2.21 }
      - { id: B, bm25: 1.2, sem: 0.7, stock_cover_days: 6, promo: 0, score: 1.58 }
      - { id: C, bm25: 0.9, sem: 0.5, stock_cover_days: 4, promo: 0, score: 1.14 }

feature_sources:
  - mongo.search_logs
  - mongo.product_features
  - mongo.products
</details>

```

## 9) Markdown Optimizer — Stochastic ladder (risk-aware), DP/MILP

### Goal & decision

Choose a **price ladder by week** to maximize **terminal gross margin** (sales GM + salvage of leftovers) by the deadline, **accounting for demand uncertainty**.

### Feature sources

- **inventory:** sku\_id , stock\_units , age\_days
- **pricing:** allowed ladder per week (e.g., {P0,P1,P2} )
- **forecast:** demand **quantiles** per (week, price)
- **products:** cogs\_cents , salvage\_cents

### Compact formula

Given stock  $S_0$ , weeks  $t = 1, \dots, T$ , ladder price  $p_t \in \mathcal{P}_t$ , margin  $m_t = p_t - \text{COGS}$ , demand  $D_t(p_t)$ , and sales  $s_t = \min(S_{t-1}, D_t(p_t))$ :

$$\max_{p_1, \dots, p_T} \mathbb{E} \left[ \sum_{t=1}^T s_t m_t + S_T \cdot \text{salvage} \right]$$

Subject to:

$$S_t = S_{t-1} - s_t, \quad 0 \leq s_t \leq S_{t-1}, \quad p_t \in \mathcal{P}_t \quad (t = 1, \dots, T).$$

**Risk-aware:** replace  $\mathbb{E}[D_t]$  by a **quantile** (e.g.,  $Q_{50}$  with a  $Q_{10}$  guard) or use **CVaR**.

## Full expansion (worked)

**Setup:**  $T = 3$ ,  $S_0 = 100$ , COGS = 6000 cents, salvage = 3000 cents.

**Ladder options:**  $P_0 = 10,000$ ,  $P_1 = 9,000$ ,  $P_2 = 8,000$  (cents)  $\Rightarrow$   
 $m(P_0) = 4000$ ,  $m(P_1) = 3000$ ,  $m(P_2) = 2000$  (cents).

**Median demand forecasts:**

$$\begin{aligned} D_1(P_0) &= 20, & D_1(P_1) &= 30, & D_1(P_2) &= 38 \\ D_2(P_0) &= 15, & D_2(P_1) &= 25, & D_2(P_2) &= 35 \\ D_3(P_0) &= 10, & D_3(P_1) &= 20, & D_3(P_2) &= 30 \end{aligned}$$

**Plan A**  $[P_0, P_1, P_2]$

$$s_1 = \min(100, 20) = 20,$$

$$GM_1 = 20 \cdot 4000 = 80,000, \quad S_1 = 100 - 20 = 80$$

$$s_2 = \min(80, 25) = 25,$$

$$GM_2 = 25 \cdot 3000 = 75,000, \quad S_2 = 80 - 25 = 55$$

$$s_3 = \min(55, 30) = 30,$$

$$GM_3 = 30 \cdot 2000 = 60,000, \quad S_3 = 55 - 30 = 25$$

$$\text{Salvage} = S_3 \cdot 3000 = 25 \cdot 3000 = 75,000$$

$$\text{Total GM}_A = 80,000 + 75,000 + 60,000 + 75,000 = \mathbf{290,000}$$

**Plan B**  $[P_1, P_1, P_1]$

$$s_1 = \min(100, 30) = 30,$$

$$GM_1 = 30 \cdot 3000 = 90,000, \quad S_1 = 100 - 30 = 70$$

$$s_2 = \min(70, 25) = 25,$$

$$GM_2 = 25 \cdot 3000 = 75,000, \quad S_2 = 70 - 25 = 45$$

$$s_3 = \min(45, 20) = 20,$$

$$GM_3 = 20 \cdot 3000 = 60,000, \quad S_3 = 45 - 20 = 25$$

$$\text{Salvage} = S_3 \cdot 3000 = 25 \cdot 3000 = 75,000$$

$$\text{Total GM}_B = 90,000 + 75,000 + 60,000 + 75,000 = \mathbf{300,000}$$

**Pick: Plan B** (higher terminal GM).

**Risk-aware note:** down-weight low-price paths if  $Q_{10}$  demand raises **stockout** risk or violates **min-margin** constraints.

▼ **metrics:** `markdown_example (with feature_sources)` — click to expand

```

metrics:
  markdown_example:
    sku_id: "SKU123"
    stock_units: 100
    cogs_cents: 6000
    salvage_cents: 3000
    ladder_prices_cents:
      W1: [10000, 9000, 8000]
      W2: [10000, 9000, 8000]
      W3: [10000, 9000, 8000]
    margins_cents: { 10000: 4000, 9000: 3000, 8000: 2000 }
    demand_med:
      W1: { 10000: 20, 9000: 30, 8000: 38 }
      W2: { 10000: 15, 9000: 25, 8000: 35 }
      W3: { 10000: 10, 9000: 20, 8000: 30 }
    plans:
      plan_A: [10000, 9000, 8000]
      plan_B: [9000, 9000, 9000]
    totals_gm_cents:
      plan_A: 290000
      plan_B: 300000

feature_sources:
  - mongo.inventory
  - mongo.pricing
  - mongo.forecast
  - mongo.products

</details>

```

## 10) Price Elasticity (GATED) — Log-log demand; Lerner rule (planning, not markdown)

### Gate to enable

Seller controls `product.price` and has  $\geq 5$  **distinct price points/SKU** (or a price-test plan). Otherwise keep **gated OFF**.

### Goal & decision

Estimate **own-price elasticity** to inform **base price**. Use with guardrails (**min margin**, **MAP**, **competitor checks**).

### Feature sources

- **price\_history**: per-SKU weekly `price_cents`, `units`
- **promos**: promo flags (control for discount effects)
- **seasonality**: week/holiday dummies
- **products**: `cogs_cents` (margins)



## Compact formula

Log-log demand (per SKU  $k$ ):

$$\ln Q_t = \alpha + \varepsilon \ln P_t + \beta^\top \mathbf{Z}_t + \epsilon_t$$

Elasticity estimate (simple OLS slope):

$$\hat{\varepsilon} = \text{SLOPE}(\ln Q, \ln P)$$

Lerner pricing (planning heuristic, if marginal cost  $MC$  known):

$$\frac{P - MC}{P} = -\frac{1}{\hat{\varepsilon}} \implies P^* = \frac{MC}{1 + 1/\hat{\varepsilon}}$$

Use with caution: need  $\hat{\varepsilon} < -1$  for an interior optimum; always respect MAP, min-margin, and competitor-band constraints.

## Full expansion (worked)

**Toy weekly data (one SKU):**  $(P, Q) = (100, 80), (110, 70), (90, 96), (95, 92), (105, 75), (115, 68)$ .

**Log transform (Excel):** =LN(P) , =LN(Q) for each row to obtain  $\ln P_t, \ln Q_t$ .

Elasticity estimate (OLS slope on logs):

$$\hat{\varepsilon} = \frac{\text{Cov}(\ln P, \ln Q)}{\text{Var}(\ln P)} \approx -1.20 \quad (\text{illustrative}).$$

**Unconstrained Lerner price (planning heuristic):** let  $MC = \text{COGS} = 60$  (same units as price).

$$P^* = \frac{MC}{1 + 1/\hat{\varepsilon}} = \frac{60}{1 + 1/(-1.20)} = \frac{60}{1 - 0.8333} = \frac{60}{0.1667} \approx 360.$$

**Guardrails (retail):** enforce **max markup band** (e.g.,  $\pm 10\%$ – $15\%$ ), **min-margin/MAP**, and **competitor match band** (e.g.,  $\pm 5\%$ ).

Use  $\hat{\varepsilon}$  to simulate candidate prices  $P$ : predict  $\hat{Q}(P)$  from the log model, then compute **revenue**  $R(P) = P \hat{Q}(P)$  and **margin**  $M(P) = (P - MC) \hat{Q}(P)$ ; choose **within guardrails**.

If  $|\hat{\varepsilon}|$  is **unstable** (few price points, promo confounding), keep **gated OFF** and gather more variation or run controlled price tests.

▼ **metrics: elasticity\_example (with feature\_sources)** — click to expand

```
metrics:
  elasticity_example:
    sku_id: "SKU789"
    cogs: 60.0
    weekly_observations:
      - { price: 100.0, units: 80 }
      - { price: 110.0, units: 70 }
      - { price: 90.0, units: 96 }
      - { price: 95.0, units: 92 }
      - { price: 105.0, units: 75 }
      - { price: 115.0, units: 68 }
    ln_series_ready: true # compute lnP, lnQ in Excel
    estimated_elasticity: -1.20 # from SLOPE(lnQ,lnP)
    lerner_unconstrained_price: 360.0
    guardrails:
      max_change_pct: 0.15
      min_margin_pct: 0.20
      map_respected: true
      competitor_match_band_pct: 0.05

feature_sources:
  - mongo.price_history
  - mongo.promos
  - mongo.products
```

</details>